


# STL and run-time speed

(C) Richèl Bilderbeek 

March 23, 2015

# Chapter 1

## Introduction

## 1.1 Overview

- Introduction
- Containers
- Algorithms
- Conclusion

## 1.2 Goal

- Do some basic container use
- Encounter some trade-offs in container choice
- Use some algorithms
- Obtain some ideas about choosing the right algorithm
- The class extensions needed to put classes in containers or use these in algorithms

# Chapter 2

# Containers

## 2.1 Question

- Name some containers.
- Distinguish between STL or non-STL, and standard or non-standard

## 2.2 Answer

- STL sequence containers:
  - `std::string`, `std::vector`, `std::list`, `std::deque`, `std::stack`
- STL associative containers:
  - `std::set`, `std::map`, `std::multi_set`, `std::multi_map`
- Standard non-STL containers:
  - `std::bitset`, `std::valarray`, `std::queue`, `std::priority_queue`
- Nonstandard non-STL containers (for C++98):
  - `std::tr1::slist`, `rope`, `hash_set`, `hash_map`
  - Boost: `array`, `circular_buffer`, `dynamic_bitset`, `graph`, `multi_array`
  - More

## 2.3 Question

- Which 'container' is native to the (C and) C++ language?
- Which STL container uses this internally?



## 2.4 Answer

- Which 'container' is native to the (C and) C++ language?
  - array
  - arrays are evil<sup>12</sup>
- Which STL container uses this internally?
  - `std::vector`
  - use `std::vector` by default<sup>3</sup>

---

<sup>1</sup>Marshall Cline, Greg Lomow and Mike Girou. C++ FAQs. ISBN: 0-201-3098301, FAQ 28.02: 'Are arrays good or evil?' (Answer: 'Arrays are evil')

<sup>2</sup>Bjarne Stroustrup. The C++ Programming Language (3rd edition). Chapter C.14.11 'Prefer vector over array'

<sup>3</sup>Herb Sutter and Andrei Alexandrescu . C++ coding standards: 101 rules, guidelines, and best practices. Chapter 76: 'Use vector by default. Otherwise choose an appropriate container'

## 2.5 Question

- Which fancy STL container is best at all of the following?
  - random access reading and writing
  - looking up elements
  - random access insertion and removal

## 2.6 Answer

- None: the characteristics are mutually exclusive
- There will be trade-offs
- Beware the illusion of container-independent code<sup>4</sup>

---

<sup>4</sup>Meyers, Effective STL, item 2

## 2.7 Question

- How to sort Persons?

```
struct Person
{
    std::string name;
    double money;
};
```

```
int main()
{
    //Sort Persons
}
```

## 2.8 Answer

- Define the global operator

```
bool operator<(const Person& lhs, const Person& rhs);
```

- Define a custom functor

```
struct SortOnMoney {  
    bool operator<(const Person& lhs, const Person& rhs) const  
};
```

## 2.9 Exercise

```
struct Person
{
    std::string name;
    double money;
};

int main()
{
    //Create some Persons
    std::vector<Person> v /* */;
    //Sort Persons on the amount of money they have
    std::sort( /* */ );
}
```

## 2.10 Question

- Describe the implementation of `std::vector`
- What are the consequences of this?

## 2.11 Answer: `std::vector`

- dynamically allocated array
- constant-time access to elements
- linear-time insertion-removal
- linear-time searching



## 2.12 Exercise: `std::vector`

- Create a class `Gossip` that prints when it is copied
- Create an empty `std::vector<Gossip>`
- Append 32 `Gossips`. How many copies are made?
- Insert a `Gossip` at the front. How many copies are made?
- Write a C-style function that works on an array of `Gossips` and swaps the first and last element

```
void SwapFirstAndLast(  
    Gossip * const gossip_array ,  
    const int size );
```

- Call `SwapFirstAndLast` and check if it did what you expected

## 2.13 Conclusion

- Use `std::vector<T>::push_back` to append to `std::vector`s, as it is amortized constant-time
- Avoid inserting elements in the front or middle of a `std::vector`
- `std::vector` can be used to communicate with C style functions

## 2.14 Question

- Describe the implementation of `std::list`
- What are the consequences of this?

## 2.15 Answer: `std::list`

- Next to the data itself, each element has two pointers: to the next and previous element in a sequential list
- constant-time insertion and removal
- linear-time access to elements
- linear-time searching

## 2.16 Exercise: `std::list`

- Create a class `Gossip` that prints when it is copied
- Create an empty `std::list<Gossip>`
- Append 32 `Gossips`. How many copies are made?
- Insert a `Gossip` at the front. How many copies are made?
- Compare `std::sort` on a `std::list` and `std::list<T>::sort`. Does one call the other?

## 2.17 Conclusion

- When adding elements to a `std::list`, no additional copies need to be made
  - Not at the end, middle, not beginning
- A `std::list` is scattered through memory
  - Calculating the number of elements is an  $O(n)$  calculation

## 2.18 Question

- Describe the implementation of `std::set`
- What are the consequences of this?

## 2.19 Answer

- Next to the data itself, each element has three pointers: to the parent, left and right branch in a red-black tree
- contents always sorted
- logarithmic-time insertion and removal
- logarithmic-time access to elements
- logarithmic-time access searching



## 2.20 Exercise

- Create a class `Person` that prints when it is copied and has at least two member variables (why two?)
- Create an empty `std::set<Person>`
- Put in 4 different `Persons`. How many copies are made?
- Create a new unique person. Use `std::set<T>::count` to check he/she is not present yet. Insert the `Person` and check he/she is present

## 2.21 Conclusion

- A `std::set` keeps its elements ordered
  - need to define operator<

# Chapter 3

# Algorithms

## 3.1 Algorithms question

- What are algorithms?
- Why use algorithms?

## 3.2 Algorithms answers

- What are algorithms?
  - named operations on multiple elements
- Why use algorithms?
  - verbosity/readability
  - algorithms are written by experts
  - algorithms are standardized: common idiom
  - increase run-time speed: naive for-loops might result in higher Big-O

### 3.3 Algorithm example

```
template<typename In, typename Out, typename Pred>
Out MysteryAlgorithm(In first, In last, Out res, Pred Pr)
{
    while (first != last)
    {
        if (Pr(*first)) *res++ = *first;
        ++first;
    }
    return res;
}
```

## 3.4 Question

- Which sorting algorithms exists?
- When to use which one?

## 3.5 Answer

- `std::sort`: when the whole range needs to be sorted
- `std::partial_sort`: when you only need the top-x values in a sorted order
- `std::nth_element`: when you only need the top-x values
- With `'stable_'` added: when the order of equivalent items needs to be preserved
- If a container has a member function with the same name, always use that one



## 3.6 Exercise 1/4

- Create a big initialized randomly-shuffled `std::vector<int>`
- Write the following functions that obtain the three lowest values:
  - `GetMinThreeUsingPartial_sort`
  - `GetMinThreeUsingNth_element`
- Check if the two functions return the same top three
- Display the top three

## 3.7 Exercise 2/4

- Create a big initialized randomly-shuffled `std::vector<int>`
- Write the following functions that obtain the three highest values:
  - `GetMaxThreeUsingPartial_sort`
  - `GetMaxThreeUsingNth_element`
  - Hint: look up `std::greater`
- Check if the two functions return the same top three
- Display the top three

## 3.8 Exercise 3/4

- Create a Person class. Every Person has two member variables:
  - a `std::string` called 'name'
  - a `double` called 'money'
- Create a big initialized randomly-shuffled `std::vector<Person>`
- Write the functions that obtain the three persons with the least money:
  - `GetMinThreeUsingPartial_sort`
  - `GetMinThreeUsingNth_element`
  - Bonus: re-use the existing ones
- Check if the two functions return the same top three
- Display the top three

## 3.9 Exercise 4/4

- Create a Person class. Every Person has two member variables:
  - a `std::string` called 'name'
  - a `double` called 'money'.
- Create a big initialized randomly-shuffled `std::vector<Person>`
- Write the functions that obtain the three persons with the most money:
  - `GetMaxThreeUsingPartial_sort`
  - `GetMaxThreeUsingNth_element`
- Check if the two functions return the same top three
- Display the top three

## 3.10 Question

- Which searching algorithms exists?
- When to use which one?

## 3.11 Answer

- `std::find`: find an element in an unsorted container
- `std::find_if`: find the first element in an unsorted container for which a predicate is true
- `std::search`: find a sequence of elements
- `std::search_n`: find an n-times-repeated sequence of elements
- `std::binary_search`: find an element in an assumed-to-be-sorted container
- `std::adjacent_find`: find two adjacent equal values
- `std::lower_bound`, `std::upper_bound`: find first/last value in a sorted container
- `std::min_element`, `std::max_element`: find min or max element

- If a container has a member function with the same name, always use that one

## 3.12 Exercise: `std::find`

- Create a Person class. Every Person has two member variables:
  - a `std::string` called 'name'
  - a `double` called 'money'
- Put some Persons in a `std::vector`, `std::list` and `std::set`
- Create a new unique person with the name 'Mr X' and put him in each of these containers
- Shuffle these containers and try to retrieve the Person with name 'Mr X'
- Create a new unique person with 123.45\$ and put him/her in each of these containers
- Shuffle these containers and try to retrieve the Person with 123.45\$



- Find the person with the most money

### 3.13 Question

- Which algorithms is used to summarize a range?

## 3.14 Answer

- `std::accumulate`
- (found in `numeric.h`)

## 3.15 Exercise

- Create a Person class. Every Person has two member variables:
  - a `std::string` called 'name'
  - a `double` called 'money'
- Put some random Persons in a container
- Obtain the sum of their money

## 3.16 Question

- What is a predicate?
- Which algorithms use a predicate?

## 3.17 Answer

- A predicate is a functor returning true or false. Identical input should yield the same results
- Algorithms:
  - `std::partition`
  - every `std::[something]_if`

## 3.18 Exercise

- Create a Person class. Every Person has two member variables:
  - a `std::string` called 'name'
  - a `double` called 'money'
- Put some random Persons in a container, of which some are poor ( $\text{money} < 1000.0$ ) and some are rich ( $\text{money} \geq 1000.0$ )
- Find a way to separate poor and rich persons in a container
- Display the poor persons
- Display the rich persons

# Chapter 4

# Conclusion



## 4.1 Conclusion

- Choose your containers with care
  - Use `std::vector` by default
- Choose your algorithms with care
  - Choose the simplest algorithm possible

# Chapter 5

EOF